

SQL Intermediate



Aggregate Functions

- COUNT, MAX, MIN, SUM, AVG
- Example:

```
SELECT max(mark)
FROM enrolment;
```

```
SELECT min(mark)
FROM enrolment;
```

```
SELECT avg(mark)
FROM enrolment;
```

```
SELECT count(stu_nbr)
FROM enrolment
WHERE mark >= 50;
```

	STU_NBR	UNIT_CODE	ENROL_YEAR	ENROL_SEMESTER	MARK	GRADE
1	11111111	FIT1001	2012	1	78	D
2	11111111	FIT1002	2013	1	(null)	(null)
3	11111111	FIT1004	2013	1	(null)	(null)
4	11111112	FIT1001	2012	1	35	N
5	11111112	FIT1001	2013	1	(null)	(null)
6	11111113	FIT1001	2012	2	65	C
7	11111113	FIT1004	2013	1	(null)	(null)
8	11111114	FIT1004	2013	1	(null)	(null)

Q1. What will be displayed by the following SQL statement?

```
SELECT count(*), count(mark)
FROM enrolment;
```

- A. 8, 8
- ☒ B. 8, 3
- C. 3, 3
- D. 3, 8

	STU_NBR	UNIT_CODE	ENROL_YEAR	ENROL_SEMESTER	MARK	GRADE
1	11111111	FIT1001	2012	1	78	D
2	11111111	FIT1002	2013	1	(null)	(null)
3	11111111	FIT1004	2013	1	(null)	(null)
4	11111112	FIT1001	2012	1	35	N
5	11111112	FIT1001	2013	1	(null)	(null)
6	11111113	FIT1001	2012	2	65	C
7	11111113	FIT1004	2013	1	(null)	(null)
8	11111114	FIT1004	2013	1	(null)	(null)

Q2. What will be displayed by the following SQL statement?

```
SELECT count(*), count(stu_nbr), count(distinct stu_nbr)
FROM enrolment;
```

- ☒ A. 8, 8, 4
- ☐ B. 8, 8, 8
- ☐ C. 8, 4, 8
- ☐ D. 8, 4, 4


	STU_NBR	UNIT_CODE	ENROL_YEAR	ENROL_SEMESTER	MARK	GRADE
1	11111111	FIT1001	2012	1	78	D
2	11111111	FIT1002	2013	1	(null)	(null)
3	11111111	FIT1004	2013	1	(null)	(null)
4	11111112	FIT1001	2012	1	35	N
5	11111112	FIT1001	2013	1	(null)	(null)
6	11111113	FIT1001	2012	2	65	C
7	11111113	FIT1004	2013	1	(null)	(null)
8	11111114	FIT1004	2013	1	(null)	(null)

Q3. We want to calculate the *average mark of the 8 rows* in the above table. What SQL statement should we use?

Note: We want to calculate $(78+35+65)/8=22.25$

- A. SELECT avg(mark) FROM enrolment; ✗
- B. SELECT sum(mark)/count(mark) FROM enrolment; ✗
- C. SELECT sum(mark)/count(*) FROM enrolment; ✓
- D. SELECT avg(NVL(mark,0)) FROM enrolment; ✓
- E. None of the above.
- F. More than one option is correct. ✓

Anatomy of an SQL Statement - Revisited



GROUP BY

- If a GROUP BY clause is used with aggregate function, the DBMS will apply the aggregate function to the different groups defined in the clause rather than all rows.

```
SELECT avg(mark)  
FROM enrolment;
```

```
SELECT unit_code, avg(mark)  
FROM enrolment  
GROUP BY unit_code  
ORDER BY unit_code;
```

```
SQL>
SQL> SELECT avg(mark)
      2  FROM enrolment;
```

```
      AVG(MARK)
-----
59.3333333
```

```
SQL>
SQL> SELECT unit_code, avg(mark)
      2  FROM enrolment
      3  GROUP BY unit_code
      4  ORDER BY unit_code;
```

```
UNIT_CO  AVG(MARK)
-----  -
FIT1001  59.3333333
FIT1002
FIT1004
```


What output is produced?

```
SELECT avg(mark)
FROM enrolmentA;
```

```
SELECT unit_code, avg(mark)
FROM enrolmentA
GROUP BY unit_code
ORDER BY unit_code;
```

```
SELECT unit_code, avg(mark), count(*)
FROM enrolmentA
GROUP BY unit_code
ORDER BY unit_code;
```

Unit_code	Mark	Studid	Year
FIT2094	80	111	2016
FIT2094	20	111	2015
FIT2004	100	111	2016
FIT2004	40	222	2015
FIT2004	40	333	2015

```
SQL> SELECT avg(mark)
2 FROM enrolmentA;
```

```
AVG(MARK)
-----
56
```

```
SQL>
SQL> SELECT unit_code, avg(mark)
2 FROM enrolmentA
3 GROUP BY unit_code
4 ORDER BY unit_code;
```

```
UNIT_CO  AVG(MARK)
-----
FIT2004   60
FIT2094   50
```

```
SQL>
SQL> SELECT unit_code, avg(mark), count(*)
2 FROM enrolmentA
3 GROUP BY unit_code
4 ORDER BY unit_code;
```

```
UNIT_CO  AVG(MARK)  COUNT(*)
-----
FIT2004   60         3
FIT2094   50         2
```

Unit_code	Mark	Studid	Year
FIT2094	80	111	2016
FIT2094	20	111	2015
FIT2004	100	111	2016
FIT2004	40	222	2015
FIT2004	40	333	2015

What output is produced?

Unit_code	Mark	Studid	Year
FIT2094	80	111	2016
FIT2094	20	111	2015
FIT2004	100	111	2016
FIT2004	40	222	2015
FIT2004	40	333	2015

```
SELECT unit_code, avg(mark), count(*)  
FROM enrolmentA  
GROUP BY unit_code, year  
ORDER BY unit_code, year;
```

```
SQL> SELECT unit_code, avg(mark), count(*)
  2  FROM enrolmentA
  3  GROUP BY unit_code, year
  4  ORDER BY unit_code, year;
```

Note: attributes in the GROUP BY clause do not have to appear in the select list

```
UNIT_CO  AVG(MARK)  COUNT(*)
-----
FIT2004      40         2
FIT2004     100         1
FIT2094      20         1
FIT2094      80         1
```

```
SQL> SELECT unit_code, year, avg(mark), count(*)
  2  FROM enrolmentA
  3  GROUP BY unit_code, year
  4  ORDER BY unit_code, year;
```

```
UNIT_CO  YEAR  AVG(MARK)  COUNT(*)
-----
FIT2004  2015      40         2
FIT2004  2016     100         1
FIT2094  2015      20         1
FIT2094  2016      80         1
```

Unit_code	Mark	Studid	Year
FIT2094	80	111	2016
FIT2094	20	111	2015
FIT2004	100	111	2016
FIT2004	40	222	2015
FIT2004	40	333	2015

HAVING clause

- It is used to put a condition or conditions on the groups defined by GROUP BY clause.

```
SELECT unit_code, count(*)  
FROM enrolment  
GROUP BY unit_code  
HAVING count(*) > 2;
```

What output is produced?

```
SELECT unit_code, avg(mark), count(*)  
FROM enrolmentA  
GROUP BY unit_code  
HAVING count(*) > 2  
ORDER BY unit_code;
```

```
SELECT unit_code, avg(mark), count(*)  
FROM enrolmentA  
GROUP BY unit_code  
HAVING avg(mark) > 55  
ORDER BY unit_code;
```

Unit_code	Mark	Studid	Year
FIT2094	80	111	2016
FIT2094	20	111	2015
FIT2004	100	111	2016
FIT2004	40	222	2015
FIT2004	40	333	2015

```
SQL> SELECT unit_code, avg(mark), count(*)
  2 FROM enrolmentA
  3 GROUP BY unit_code
  4 HAVING count(*) > 2
  5 ORDER BY unit_code;
```

```
UNIT_CO  AVG(MARK)  COUNT(*)
-----  -
FIT2004      60         3
```

```
SQL>
SQL> SELECT unit_code, avg(mark), count(*)
  2 FROM enrolmentA
  3 GROUP BY unit_code
  4 HAVING avg(mark) > 55
  5 ORDER BY unit_code;
```

```
UNIT_CO  AVG(MARK)  COUNT(*)
-----  -
FIT2004      60         3
```

Unit_code	Mark	Studid	Year
FIT2094	80	111	2016
FIT2094	20	111	2015
FIT2004	100	111	2016
FIT2004	40	222	2015
FIT2004	40	333	2015

HAVING and WHERE clauses

```
SELECT unit_code, count(*)  
FROM enrolment  
WHERE mark IS NULL  
GROUP BY unit_code  
HAVING count(*) > 1;
```

- The WHERE clause is applied to ALL rows in the table.
- The HAVING clause is applied to the groups defined by the GROUP BY clause.
- The order of operations performed is FROM, WHERE, GROUP BY, HAVING and then ORDER BY.
- On the above example, the logic of the process will be:
 - All rows where mark is NULL are retrieved. (due to the WHERE clause)
 - The retrieved rows then are grouped into different unit_code.
 - If the number of rows in a group is greater than 1, the unit_code and the total is displayed. (due to the HAVING clause)

What output is produced?

Unit_code	Mark	Studid	Year
FIT2094	80	111	2016
FIT2094	20	111	2015
FIT2004	100	111	2016
FIT2004	40	222	2015
FIT2004	40	333	2015

```
SELECT unit_code, avg(mark), count(*)  
FROM enrolmentA  
WHERE year = 2015  
GROUP BY unit_code  
HAVING avg(mark) > 30  
ORDER BY avg(mark) DESC;
```

```

SQL> SELECT unit_code, avg(mark), count(*)
  2  FROM enrolmentA
  3  WHERE year = 2015
  4  GROUP BY unit_code
  5  HAVING avg(mark) > 30
  6  ORDER BY avg(mark) DESC;

```

```

UNIT_CO  AVG(MARK)  COUNT(*)
-----  -
FIT2004      40         2

```

Unit_code	Mark	Studid	Year
FIT2094	80	111	2016
FIT2094	20	111	2015
FIT2004	100	111	2016
FIT2004	40	222	2015
FIT2004	40	333	2015

Unit_code	Mark	Studid	Year
FIT2094	80	111	2016
FIT2094	20	111	2015
FIT2004	100	111	2016
FIT2004	40	222	2015
FIT2004	40	333	2015

Q4. What is the output for:

```
SELECT unit_code, studid, avg(mark)
FROM enrolmentA
GROUP BY unit_code
HAVING avg(mark) > 55
ORDER BY unit_code, studid;
```

- A. FIT2094, 111, 50
- B. FIT2004, 111, 60
- C. FIT2004, 111, 60, 222, 333
- D. FIT2004, 111, 100
- E. Will print three rows
- F. Error

```
SQL> SELECT unit_code, studid, avg(mark)
  2  FROM enrolmentA
  3  GROUP BY unit_code
  4  HAVING avg(mark) > 55
  5  ORDER BY unit_code, studid;
```

Error starting at line : 1 in command -
 SELECT unit_code, studid, avg(mark)
 FROM enrolmentA
 GROUP BY unit_code
 HAVING avg(mark) > 55
 ORDER BY unit_code, studid
 Error at Command Line : 1 Column : 19
 Error report -
 SQL Error: ORA-00979: **not a GROUP BY expression**
 00979. 00000 - "not a GROUP BY expression"
 *Cause:
 *Action:

Unit_code	Mark	Studid	Year
FIT2094	80	111	2016
FIT2094	20	111	2015
FIT2004	100	111	2016
FIT2004	40	222	2015
FIT2004	40	333	2015

```
SELECT stu_lname, stu_fname, avg(mark)
FROM enrolment e JOIN student s
      ON s.stu_nbr = e.stu_nbr
GROUP BY s.stu_nbr;
```

The above SQL generates error message

```
SQL Error: ORA-00979: not a GROUP BY expression
00979. 00000 - "not a GROUP BY expression"
```

Why and how to fix this?

- Why? Because the grouping is based on the stu_nbr, whereas the display is based on stu_lname and stu_fname. The two groups may not have the same members.
- How to fix this?
 - Include the stu_lname,stu_fname as part of the GROUP BY condition.
- Attributes that are used in the SELECT, HAVING and ORDER BY must be included in the GROUP BY clause.

Subqueries

- Query within a query.

"Find all students whose mark is higher than the average mark of all enrolled students"

```
SELECT *  
FROM enrolment  
WHERE mark > (SELECT avg (mark)  
               FROM enrolment );
```

Types of Subqueries


Single-value



Multiple-row subquery (a list of values – many rows, one column)



Multiple-column subquery (many rows, many columns)



Q5. What will be returned by the *inner query*?

```
SELECT *  
FROM enrolment  
WHERE mark > (SELECT avg(mark)  
              FROM enrolment  
              GROUP BY unit_code);
```

- A. A value (a single column, single row).
- ☒ B. A list of values.
- C. Multiple columns, multiple rows.
- D. None of the above.


```
SQL> SELECT *  
  2  FROM enrolment  
  3  WHERE mark > (SELECT avg(mark)  
  4                FROM enrolment  
  5                GROUP BY unit_code);
```

Error starting at line : 1 in command -

SELECT *

FROM enrolment

WHERE mark > (SELECT avg(mark)

FROM enrolment

GROUP BY unit_code)

Error report -

**ORA-01427: single-row subquery returns more than one
row**

Q6. What will be returned by the *inner query*?

```
SELECT unit_code, stu_lname, stu_fname, mark
FROM enrolment e join student s
    on e.stu_nbr = s.stu_nbr
WHERE (unit_code, mark) IN (SELECT unit_code, max(mark)
    FROM enrolment
    GROUP BY unit_code);
```

- A. A value (a single column, single row).
- B. A list of values.
- ☒ C. Multiple columns, multiple rows.
- D. None of the above.

Comparison Operators for Subquery

- Operator for single value comparison.
=, <, >
- Operator for multiple rows or a list comparison.
 - equality
 - IN
 - inequality
 - ALL, ANY combined with <, >

	STU_NBR	UNIT_CODE	ENROL_YEAR	ENROL_SEMESTER	MARK	GRADE
1	11111111	FIT1001	2012	1	78	D
2	11111111	FIT1002	2013	1	80	HD
3	11111111	FIT1004	2013	1	85	HD
4	11111112	FIT1001	2012	1	35	N
5	11111112	FIT1001	2013	1	50	P
6	11111113	FIT1001	2012	2	65	C
7	11111113	FIT1004	2013	1	89	HD
8	11111114	FIT1004	2013	1	50	P

Q7. Which row(s) in ENROL2 table will be retrieved by the following SQL statement?

```
SELECT * FROM enrol2
WHERE mark IN (SELECT max(mark)
               FROM enrol2
               GROUP BY unit_code);
```

- ☒ A. 1, 2, 7
- ☐ B. 7
- ☐ C. 2, 3, 7

	STU_NBR	UNIT_CODE	ENROL_YEAR	ENROL_SEMESTER	MARK	GRADE
1	11111111	FIT1001	2012	1	78	D
2	11111111	FIT1002	2013	1	80	HD
3	11111111	FIT1004	2013	1	85	HD
4	11111112	FIT1001	2012	1	35	N
5	11111112	FIT1001	2013	1	50	P
6	11111113	FIT1001	2012	2	65	C
7	11111113	FIT1004	2013	1	89	HD
8	11111114	FIT1004	2013	1	50	P

```
SQL> SELECT * FROM enrol2
2 WHERE mark IN (SELECT max(mark)
3                 FROM enrol2
4                 GROUP BY unit_code)
5 ORDER BY stu_nbr, unit_code, enrol_year;
```

STU_NBR	UNIT_CO	ENROL_YEAR	E	MARK	GRA
11111111	FIT1001	2012	1	78	D
11111111	FIT1002	2013	1	80	HD
11111113	FIT1004	2013	1	89	HD

	STU_NBR	UNIT_CODE	ENROL_YEAR	ENROL_SEMESTER	MARK	GRADE
1	11111111	FIT1001	2012	1	78	D
2	11111111	FIT1002	2013	1	80	HD
3	11111111	FIT1004	2013	1	85	HD
4	11111112	FIT1001	2012	1	35	N
5	11111112	FIT1001	2013	1	50	P
6	11111113	FIT1001	2012	2	65	C
7	11111113	FIT1004	2013	1	89	HD
8	11111114	FIT1004	2013	1	50	P

UCODE	ROUND(AVG(MARK))
FIT1001	57
FIT1002	80
FIT1004	75

Q8. Which row/s in ENROL2 will be retrieved by the following SQL statement?

SELECT * FROM enrol2

WHERE mark > ANY (SELECT avg(mark)

FROM enrol2

GROUP BY unit_code);

see smallest

- A. 1, 2, 3, 6, 7
 B. 2, 3, 7
 C. 3, 7
 D. No rows will be returned

	STU_NBR	UNIT_CODE	ENROL_YEAR	ENROL_SEMESTER	MARK	GRADE
1	11111111	FIT1001	2012	1	78	D
2	11111111	FIT1002	2013	1	80	HD
3	11111111	FIT1004	2013	1	85	HD
4	11111112	FIT1001	2012	1	35	N
5	11111112	FIT1001	2013	1	50	P
6	11111113	FIT1001	2012	2	65	C
7	11111113	FIT1004	2013	1	89	HD
8	11111114	FIT1004	2013	1	50	P

UCODE	ROUND(AVG(MARK))
FIT1001	57
FIT1002	80
FIT1004	75

```
SQL> SELECT * FROM enrol2
      2 WHERE mark > ANY (SELECT avg(mark)
      3                      FROM enrol2
      4                      GROUP BY unit_code)
      5 ORDER BY stu_nbr, unit_code, enrol_year, enrol_semester;
```

STU_NBR	UNIT_CO	ENROL_YEAR	E	MARK	GRA
11111111	FIT1001	2012	1	78	D
11111111	FIT1002	2013	1	80	HD
11111111	FIT1004	2013	1	85	HD
11111113	FIT1001	2012	2	65	C
11111113	FIT1004	2013	1	89	HD

	STU_NBR	UNIT_CODE	ENROL_YEAR	ENROL_SEMESTER	MARK	GRADE
1	11111111	FIT1001	2012	1	78	D
2	11111111	FIT1002	2013	1	80	HD
3	11111111	FIT1004	2013	1	85	HD
4	11111112	FIT1001	2012	1	35	N
5	11111112	FIT1001	2013	1	50	P
6	11111113	FIT1001	2012	2	65	C
7	11111113	FIT1004	2013	1	89	HD
8	11111114	FIT1004	2013	1	50	P

R1	UCODE	R2	ROUND(AVG(MARK))
	FIT1001		57
	FIT1002		80
	FIT1004		75

Q9. Which row/s in ENROL2 will be retrieved by the following SQL statement?

SELECT * FROM enrol2
WHERE mark > ALL (SELECT avg(mark)
FROM enrol2
GROUP BY unit_code);

see largest

- A. 1, 2, 3, 6, 7
- B. 2, 3, 7
- C. 3, 7
- D. No rows will be returned

	STU_NBR	UNIT_CODE	ENROL_YEAR	ENROL_SEMESTER	MARK	GRADE
1	11111111	FIT1001	2012	1	78	D
2	11111111	FIT1002	2013	1	80	HD
3	11111111	FIT1004	2013	1	85	HD
4	11111112	FIT1001	2012	1	35	N
5	11111112	FIT1001	2013	1	50	P
6	11111113	FIT1001	2012	2	65	C
7	11111113	FIT1004	2013	1	89	HD
8	11111114	FIT1004	2013	1	50	P

UCODE	ROUND(AVG(MARK))
FIT1001	57
FIT1002	80
FIT1004	75

```
SQL> SELECT * FROM enrol2
      2 WHERE mark > ALL (SELECT avg(mark)
      3                      FROM enrol2
      4                      GROUP BY unit_code)
      5 ORDER BY stu_nbr, unit_code, enrol_year, enrol_semester;
```

STU_NBR	UNIT_CO	ENROL_YEAR	E	MARK	GRA
11111111	FIT1004	2013	1	85	HD
11111113	FIT1004	2013	1	89	HD

Q10. Find all students whose mark in any enrolled unit is lower than Wendy Wheat's lowest mark for all units she is enrolled in. What would be a possible inner query statement for the above query (assume Wendy Wheat's name is unique)?

A. SELECT min(mark)

FROM enrol2

WHERE stu_lname='Wheat' AND stu_fname='Wendy';



B. SELECT min(mark)

FROM enrol2 e JOIN student s on e.studid = s.studid

WHERE stu_lname='Wheat' AND stu_fname='Wendy';



C. SELECT min(mark) FROM enrol2;



D. SELECT mark

FROM enrol2 e JOIN student s on e.studid = s.studid

WHERE stu_lname='Wheat' AND stu_fname='Wendy';



Summary

- Aggregate Functions
 - count, min, max, avg, sum
- GROUP BY and HAVING clauses.
- Subquery
 - Inner vs outer query
 - comparison operators (IN, ANY, ALL)